

Entering / cleaning data 1

Odds and ends

- Next week, we will not meet for either class session (Aug. 27 and 29). Instead, videos of the lectures will be available through the online coursebook. You will be responsible for watching these sometime during next week on your own time.
- Office hours will be in EH 120, 10:00–11:00 AM Fridays. For the office hours on Aug. 31, we will go through the group exercises for Week 1 (“Entering and Cleaning Data #1”). Attendance is voluntary, but you are responsible for trying the exercises whether you attend this in-person session or not.

- The first quiz will be on Wednesday, Sept. 5 (next in-person class meeting). It will cover the content in Chapter 1 of the online book (“R Preliminaries”). The vocabulary list is up in the “Vocabulary” appendix of the online book.
- The first homework is due Sept. 12. The updated assignment for 2018 is available now in the online coursebook.

Getting data into R

Basics of getting data into R

Basic approach:

- Download data to your computer
- Save the data in your R Project directory for the project you're using it for or in a subdirectory within that directory ("data" is a good name for this subdirectory)
- Read data into R (functions in readr: `read_csv`, `read_table`, `read_delim`, `read_fwf`, etc.)
- Check to make sure the data came in correctly (`dim`, `ncol`, `nrow`, `head`, `tail`, `str`, `colnames`)

Reading data into R

What kind of data can you get into R?

The sky is the limit. . .

- **Flat files**
- Files from other statistical packages (SAS, Excel, Stata, SPSS)
- Tables on webpages (e.g., the table near the end of this page)
- Data in a database (e.g., SQL)
- Data stored in XML and JSON
- Really crazy data formats used in other disciplines (e.g., netCDF files from climate folks, MRI data stored in Analyze, NIfTI, and DICOM formats)
- Data through APIs (e.g., GoogleMaps, Twitter, many government agencies)
- Incredibly messy data using `scan` and `readLines`

Flat files

R can read in data from *a lot* of different formats. The only catch: you need to tell R how to do it.

To start, we'll look at **flat files**, which are plain text files (i.e., you can read them when you open them in a text editor, unlike a file in a binary format, like an Excel or Word file) with a two-dimensional structure (a row for each observation and a column for each variable).

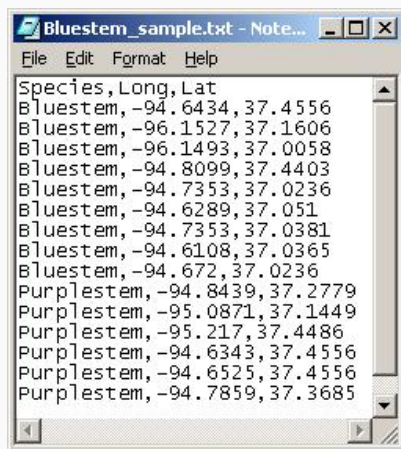
Types of flat files

There are two main types of flat files:

1. **Fixed width files:** Each column is a certain number of characters wide. (If you printed it out, you could draw vertical lines that separate the columns.)
2. **Delimited files:** In each row, a certain symbol (**delimiter**) separates the data into columns values for that observation.
 - “.csv”: Comma-separated values
 - “.tab”, “.tsv”: Tab-separated values
 - Other possible delimiters: colon, semicolon, pipe (“|”)

See if you can identify what types of files the following files are. . .

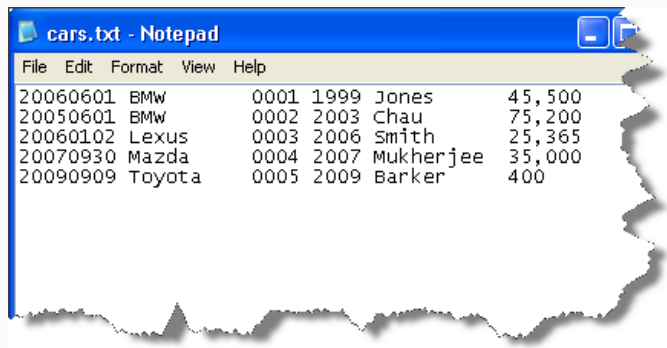
What type of file?



A screenshot of a Notepad window titled "Bluestem_sample.txt - Note...". The window has a menu bar with "File", "Edit", "Format", and "Help". The text inside the window is as follows:

```
Species,Long,Lat
Bluestem,-94.6434,37.4556
Bluestem,-96.1527,37.1606
Bluestem,-96.1493,37.0058
Bluestem,-94.8099,37.4403
Bluestem,-94.7353,37.0236
Bluestem,-94.6289,37.051
Bluestem,-94.7353,37.0381
Bluestem,-94.6108,37.0365
Bluestem,-94.672,37.0236
Purplestem,-94.8439,37.2779
Purplestem,-95.0871,37.1449
Purplestem,-95.217,37.4486
Purplestem,-94.6343,37.4556
Purplestem,-94.6525,37.4556
Purplestem,-94.7859,37.3685
```

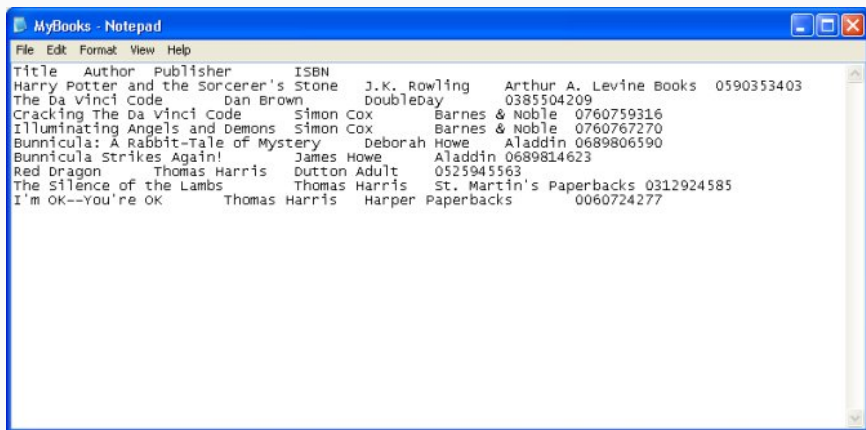
What type of file?



What type of file?

```
H|20110606|pizza.txt|
D|10|Chicken Pesto|20|23|30|5.5|7.4|9.9||
D|10|Meatball|10|53|60|6.5|8.4|10.9|
D|10|Fire Cracker|3|13|60|5.8|7.9|11.9|
D|10|Spinach|1|2|5|5.5|7.0|8.8|
D|10|BBQ Chicken|35|102|95|6.5|7.9|10.9|
D|10|Vegetarian|5|13|28|4.5|7.9|9.5|
D|10|Mexican|11|33|36|5.5|7.4|9.9|
D|10|The Monaco|22|53|7|5.5|7.5|8.9|
D|10|Chilli Prawn|5|5|6|5.5|7.4|9.9|
D|10|Chefs Special|8|18|40|5.8|7.8|9.8|
D|10|Marinara|3|17|41|5.5|7.4|9.0|
D|10|Supreme|50|52|58|5.5|7.4|9.2|
D|10|Margherita|9|19|87|5.0|7.0|8.0|
D|10|Napoli|60|85|66|5.2|7.2|9.2|
D|10|Caprice|31|32|38|5.5|7.4|9.3|
D|10|Ham and Pineapple|18|39|28|5.8|7.0|9.0|
T|16|
```

What type of file?



What type of file?

File Edit Format View Help

Title, Subtitle, Larger work, Contributor #1, Contributor #2, Contributor #3, Contributor #4, Genre, Publisher, Published Location, Date
Published, Instrumentation, Key, Location, Indiana Connection, Sheet Music
Consortium, Notes, Complete
""A"" "You're Adorable", The alphabet song,, Buddy Kaye, Sidney Lippman, Fred Wise,, Popular standard, Laurel Music Corporation, "New York, NY", 1948, Voice and piano/guitar or ukulele, C Major,, None, Yes, Perry Como pictured on cover,
"Aba Daba Honey Moon, The",,, ""Two weeks with Love"" Motion Picture", Arthur Fields, Walter Donovan,, "Popular Standard, Movie Selection", Leo Feist Inc., "New York, NY", 1942, Voice and Piano, C Minor,, None, Yes,,
Abi Bezant,, ""Mamele"" Motion Picture", Abraham Ellstein, Molly Picon,, "Popular Standard, Movie Selection", Metro Music Co., "New York, NY", 1939, Voice and Piano, E Minor,, None, No, Molly Picon pictured on cover,
Abdul the Bulbul Ameer,, Bob Kaai, Jim Smock,, Popular Standard, Calumet Music Co., "Chicago, IL", 1935, "Voice, Piano, Hawaiian Guitar, Ukulele", G Major,, None, Yes, Ben Pollack pictured on cover,
About A Quarter to Nine,, ""Go Into Your Dance"" Motion Picture", Harry Warren, Al Dubin,, "Popular Standard, Movie Selection", M. Witmark & Sons, "New York, NY", 1935, "Voice, Piano, Guitar, Ukelele", E Minor,, None, No, Al Jolson and Ruby Keeler pictured on cover,
Absent,, John. W. Metcalf, Catherine Young Glen,, Popular Standard, Arthur P. Schmidt, "Boston, MA", 1899, Voice and Piano, G Major,, None, Yes,,
The Academy Two-Step,, , Barclay Walker,, , Popular Standard, Carlin & Lennox, "Indianapolis, IN", , Piano, F Major,, Composer, No,,
Ac-cent-tchu-ate the Positive, Mister In Between,, ""Here Come the Waves"" Motion Picture", Harold Arlen, Johnny Mercer,, "Popular Standard, Movie Selection", Edwin H. Morris & Co., "New York, NY", 1944, "Voice, Piano, Guitar", F Major,, None, Yes, Bing Crosby and Betty Hutton pictured on cover,
Across the Alley From the Alamo,, , Joe Greene,, , Popular Standard, Leslie Music

What type of file?

1000233	Miralda	John
1000234	Faley	Nick
1000235	Baylog	Cathy
1000236	Gallardo	Mike
1000237	Christian	Daniel
1000238	Baufield	Daniel
1000239	Frazier	Robert
1000240	Garrido	Edward
1000241	Williams	Zachary
1000242	Morel	David
	Padilla	Damian
1000244	Rosenberg	Wayne
1000245	Blanchard	Phong S
1000246	Wiggins	David
1000247	Miller	Jeffrey
1000248	Coon	Terry
1000249	Chretien	Walter
1000250	Myers	Timothy
1000233	Miralda	John
1000234	Faley	Nick
1000235	Baylog	Cathy

Types of flat files

Flat files will often end in file extensions like “.txt”, “.csv”, “.fwf”, and “.tsv”.

To figure out the structure of a flat file, start by opening it in a text editor. RStudio can also be used as a text editor to open and explore flat files (right click on the file name and then choose “Open With” and “RStudio”).

Reading in flat files

R can read any of the types of files we just looked at by using one of the functions from the `readr` package:

File type	General function
Delimited	<code>read_delim</code>
Fixed width	<code>read_fwf</code>

You will just need to be able to clearly tell R *how* to read the file in, including what type of flat file it is and what delimiter it uses.

Reading in flat files

For example, the file “AWOIS_Wrecks_KnownYear.tab” is a flat delimited file with tabs as delimiters containing the subset of the Office of the Coast Survey’s Automated Wreck and Obstruction Information System (AWOIS) for which the year the vessel sank is known.

You can download this file by going to this link and using the “Raw” button in the top right hand corner (right click and select “Download Linked File”).

Reading in flat files

If save this file in your working directory, to read it in and assign it the name `shipwrecks`, you can run:

```
library(readr)
shipwrecks <- read_delim("AWOIS_Wrecks_KnownYear.tab",
                        delim = "\t")
```

Check out a subset of the data

```
shipwrecks[c(1:4), c(2, 4, 5, 9)]
```

```
## # A tibble: 4 x 4
##   VESSLTERMS    LATDEC LONDEC YEARSUNK
##   <chr>        <dbl>  <dbl>    <int>
## 1 SUBCHASER 187    37.3   -75.5    1918
## 2 BIRCH LAKE    37.3   -75.6    1943
## 3 PACIFIC      37.3   -75.6    1925
## 4 UNKNOWN      37.3   -75.7    1916
```

readr family of functions

Some of the interesting options with the readr family of functions are:

Option	Description
<code>skip</code>	How many lines of the start of the file should you skip?
<code>col_names</code>	What would you like to use as the column names?
<code>col_types</code>	What would you like to use as the column types?
<code>n_max</code>	How many rows do you want to read in?
<code>na</code>	How are missing values coded?

Reading in flat files

The “daily_show_guests.csv” file you worked with in the previous In-Course Exercise is a delimited flat file with commas as the delimiters. **It also has four lines of information about the data, before the actual data begins.**

```
1  ## Obtained from GitHub page of FiveThirtyEight under the
2  ## Creative Commons Attribution 4.0 International License
3  ## https://github.com/fivethirtyeight/data/tree/master/daily-show-guests
4  ##
5  YEAR,GoogleKnowlege_Occupation,Show,Group,Raw_Guest_List
6  1999,actor,1/11/99,Acting,Michael J. Fox
7  1999,Comedian,1/12/99,Comedy,Sandra Bernhard
8  1999,television actress,1/13/99,Acting,Tracey Ullman
```

Reading in flat files

You can handle this by using the skip option to tell R to skip the first four lines:

```
read_delim("daily_show_guests.csv", delim = ",", skip = 4)
```

```
## # A tibble: 2,693 x 5
##   YEAR GoogleKnowlege_Occupati~ Show    Group
##   <int> <chr>                  <chr>   <chr>
## 1  1999 actor                  1/11/99 Acting
## 2  1999 Comedian              1/12/99 Comedy
## 3  1999 television actress    1/13/99 Acting
## 4  1999 film actress          1/14/99 Acting
## 5  1999 actor                  1/18/99 Acting
## 6  1999 actor                  1/19/99 Acting
## 7  1999 Singer-lyricist        1/20/99 Musici~
## 8  1999 model                  1/21/99 Media
## 9  1999 actor                  1/25/99 Acting
## 10 1999 stand-up comedian      1/26/99 Comedy
## # ... with 2,683 more rows, and 1 more variable:
## #   Raw_Guest_List <chr>
```

readr family of functions

Many members of the `readr` package that read delimited files are doing the same basic thing. The only difference is what defaults they have for the separator (`delim`).

Some key members of the `readr` family for delimited data:

Function	Separator
<code>read_csv</code>	comma
<code>read_csv2</code>	semi-colon
<code>read_table2</code>	whitespace
<code>read_tsv</code>	tab

readr family of functions

For any type of delimited flat files, you can also use the more general `read_delim` function to read in the file. However, you will have to specify yourself what the delimiter is (e.g., `delim = ","` for a comma-separated file).

For example, the following two calls do the same thing:

```
read_delim("daily_show_guests.csv", delim = ",", skip = 4)
read_csv("daily_show_guests.csv", skip = 4)
```


The readr package also includes some functions for reading in fixed width files:

- `read_fwf`
- `read_table`

These allow you to specify field widths for each fixed width field, but they will also try to determine the field-widths automatically.

Reading data from other files types

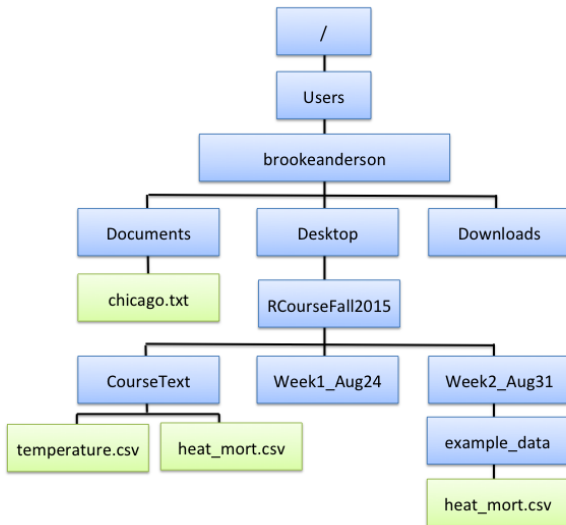
You can also read data in from a variety of other file formats, including:

File type	Function	Package
Excel	<code>read_excel</code>	<code>readxl</code>
SAS	<code>read_sas</code>	<code>haven</code>
SPSS	<code>read_spss</code>	<code>haven</code>
Stata	<code>read_stata</code>	<code>haven</code>

We'll take another break here to work on the next part of the In-Course Exercise (Section 2.7.1 and 2.7.2).

Directory structure

Computer directory structure



Working directory

The **working directory** is the directory within your directory structure from which your R session is currently working.

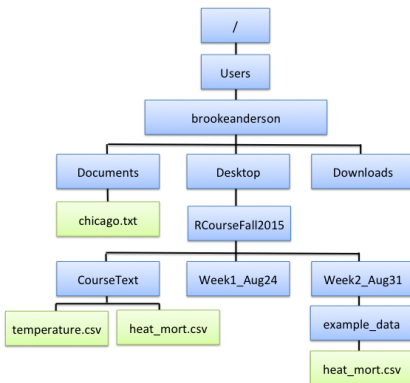
When you use **R Projects** to organize your work and files in R, anytime you open one of your R Projects, your working directory will automatically be that project's directory.

Working directory

To confirm this, open one of your R Projects and print out your working directory using the function `getwd()`:

```
getwd()
```

```
[1] "/Users/brookeanderson/Desktop/RCourseFall2015"
```

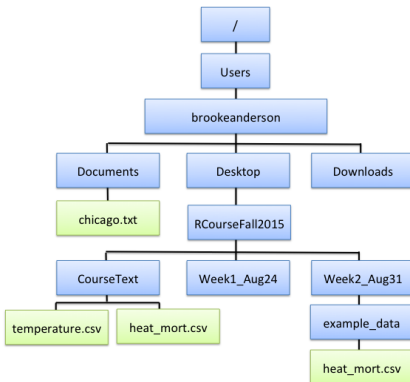


Working directory

You can use the `list.files()` function to list all the files in your current working directory:

```
list.files()
```

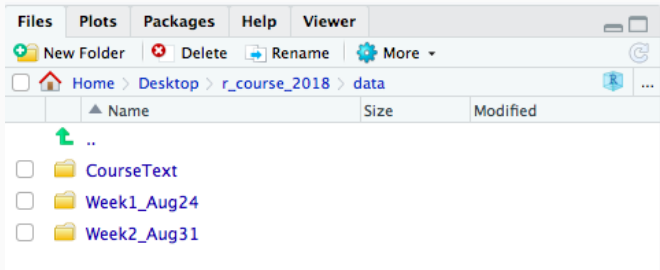
```
[1] "CourseText"      "Week1_Aug24"     "Week2_Aug31"
```



Working directory

The “Files” pane in RStudio (often on the lower right) will also show you the files available in your current working directory.

This should line up with what you get if you run `list.files()`.

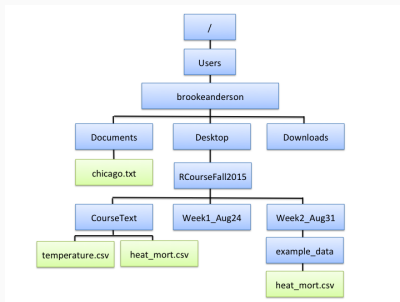


Working directory

When you run `list.files()`, if there is a name without a file extension, it's probably the name of a **subdirectory** of your current working directory. To list the files in one of these subdirectories, you can use `list.files` with that subdirectory's name.

```
list.files("CourseText")
```

```
[1] "temperature.csv"      "heat_mort.csv"
```



Taking advantage of paste0

You can create an object with your directory name using `paste0`, and then use that to set your directory. We'll take a lot of advantage of this for reading in files.

The convention for `paste0` is:

```
## Generic code
[object name] <- paste0("[first thing you want to paste]",
                        "[what you want to add to that]",
                        "[more you want to add]")
```

Taking advantage of paste0

Here's an example:

```
my_dir <- paste0("~/RProgrammingForResearch",  
                "data/measles_data")
```

```
my_dir
```

```
## [1] "~/RProgrammingForResearchdata/measles_data"
```

```
setwd(my_dir)
```

Relative versus absolute pathnames

When you want to reference a directory or file that is not in your working directory, you need to give R the directions for how to find the file. You can use one of two types of pathnames:

- *Relative pathname*: How to get to the file or directory from your current working directory
- *Absolute pathname*: How to get to the file or directory from anywhere on the computer

Relative versus absolute pathnames

If directions worked like pathnames, here is how you would tell someone to get to this building:

- *Relative pathname*: Turn right on Center, then turn right when you get to the intersection, then go to the second building on your left (only works if you started on Prospect going west right before the intersection with Center).
- *Absolute pathname*: Go to 350 W. Lake Street, Fort Collins, CO, USA.

Relative versus absolute pathnames

Say your current working directory was `/Users/brookeanderson/RProgrammingForResearch` and you wanted to get into the subdirectory `data`. Here are examples of referencing that subdirectory using the two types of pathnames:

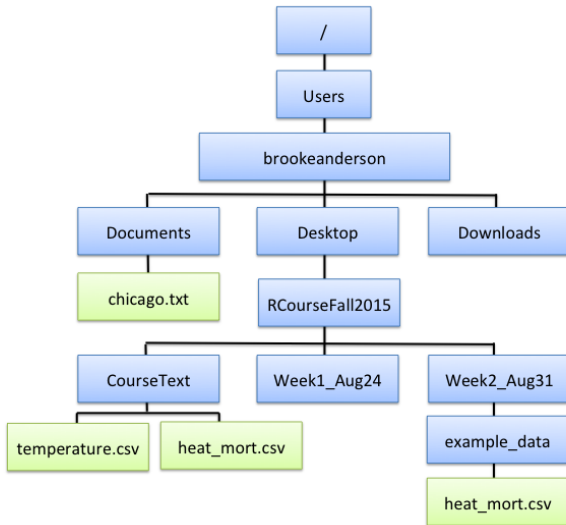
Absolute:

```
"/Users/brookeanderson/RProgrammingForResearch/data"
```

Relative:

```
"data"
```

Computer directory structure



Getting around directories

There are a few abbreviations you can use to represent certain relative locations when you're using `setwd()`:

Shorthand	Meaning
.	Current working directory
..	One directory up from current working directory (parent directory)
../..	Two directories up from current working directory
../data	The 'data' subdirectory of the parent directory

Relative versus absolute pathnames

Here are some examples of relative pathnames that use these abbreviations:

If data is a subdirectory of your current parent directory (i.e., from your working directory, you need to go “up and over”):

```
"../data"
```

If data is a subdirectory of the subdirectory Ex of your current working directory:

```
"Ex/data"
```

Relative versus absolute pathnames

Both methods of writing filenames have their own advantages and disadvantages:

- *Relative pathname*: Which file you are indicating depends on which working directory you are in, which means that your code will break if you try to re-run it from a different working directory. However, relative pathways in your code make it easier for you to share a working version of a project with someone else. For most of this course, we will focus on using relative pathnames, especially when you start collaborating.
- *Absolute pathname*: No matter what working directory you're in, it is completely clear to your computer which file you mean when you use an absolute pathname. However, your code will not work on someone else's computer without modifications (because the structure of their computer's full directory will be different).

Relative versus absolute pathnames

I **strongly** recommend saving your data files somewhere in the directory structure of the R Project in which you're working and then using **relative pathnames** to reference that file when you need to read it in.

This practice establishes good habits for making your research computationally reproducible.

We'll take a break now to do the first part of the in-course exercise (Sections 2.7.3).

Tidyverse

The “tidyverse”

You will also see code that uses functions like `read.csv` and `read.table` to read in flat files. These are from the `read.table` family of functions, which are part of base R. The `readr` functions are very similar, but have some more sensible defaults, including in determining column classes.

Compared to the `read.table` family of functions, the `read_*` functions:

- Work better with large datasets: faster, includes progress bar
- Have more sensible defaults (e.g., characters default to characters, not factors)

The “tidyverse”

The readr package is part of the “tidyverse”— a collection of recent and developing packages for R, many written by Hadley Wickham.



The “tidyverse”



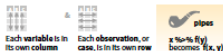
"A giant among data nerds"

<https://priceonomics.com/hadley-wickham-the-man-who-revolutionized-r/>

Data Transformation with dplyr : CHEAT SHEET




dplyr functions work with pipes and expect tidy data. In tidy data:



Summarise Cases

These apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).



VARIATIONS

- summarise_all()** - Apply funs to every column.
- summarise_at()** - Apply funs to specific columns.
- summarise_if()** - Apply funs to all cols of one type.

Group Cases

Use **group_by()** to create a "grouped" copy of a table. dplyr functions will manipulate each "group" separately and then combine the results.




group_by(data, ..., add = FALSE)
Returns copy of table grouped by ...
giris <- group_by(mtcars, Species)

ungroup(x, ...)
Returns ungrouped copy of table.
ungroup(giris)

Manipulate Cases

EXTRACT CASES

Row functions return a subset of rows as a new table.



Logical and boolean operators to use with filter()

```
<      <=      is.na()  %in%      |      xor()
>      >=      !is.na()  !        &
```

See ?base::logic and ?Comparison for help.

ARRANGE CASES

arrange(data, ...) Order rows by values of a column or columns (low to high), use with **desc()** to order from high to low.
arrange(mtcars, mpg)
arrange(mtcars, desc(mpg))


ADD CASES

add_row(data, ..., before = NULL, after = NULL)
Add one or more rows to a table.
add_row(kothifur, eruptions = 1, waiting = 1)

Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.




Use these helpers with **select()**, e.g. **select(mtcars, starts_with("Sepsol"))**

contains(match) **num_range(prefix, range)** ; e.g. **mpg:cyl**
ends_with(match) **one_of(...)** ; e.g. **Species**
matches(match) **starts_with(match)**

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).

vectorized function



RStudio has several very helpful **cheatsheets**. These are one-page sheets (front and back) that cover many of the main functions for a certain topic or task in R. These cheatsheets cover a lot of the main “tidyverse” functions.

You can access these directly from RStudio. Go to “Help” -> “Cheatsheets” and select the cheatsheet on the topic of interest.

You can find even more of these cheatsheets at <https://www.rstudio.com/resources/cheatsheets/>.

Data cleaning

Cleaning data

Common data-cleaning tasks include:

Task	dplyr function
Renaming columns	<code>rename</code>
Selecting certain columns	<code>select</code>
Adding or changing columns	<code>mutate</code>
Limiting to certain rows	<code>slice</code>
Filtering to certain rows	<code>filter</code>
Arranging rows	<code>arrange</code>

The “tidyverse”

Today, we'll talk about using functions from the `dplyr` and `lubridate` packages, which are both part of the “tidyverse”, like the `readr` package.

To use these functions, you'll need to load those packages:

```
library("dplyr")  
library("lubridate")
```

```
## Warning: package 'lubridate' was built under R  
## version 3.4.4
```

Cleaning data

As an example, let's look at the Daily Show data:

```
daily_show <- read_csv("../data/daily_show_guests.csv",  
                        skip = 4)  
head(daily_show, 3)
```

```
## # A tibble: 3 x 5  
##   YEAR GoogleKnowlege_Occupation Show      Group  
##   <int> <chr>                  <chr>    <chr>  
## 1  1999 actor                  1/11/99 Acting  
## 2  1999 Comedian              1/12/99 Comedy  
## 3  1999 television actress    1/13/99 Acting  
## # ... with 1 more variable: Raw_Guest_List <chr>
```

Re-naming columns

A first step is often re-naming columns. It can be hard to work with a column name that is:

- long
- includes spaces
- includes upper case

Several of the column names in `daily_show` have some of these issues:

```
colnames(daily_show)
```

```
## [1] "YEAR"  
## [2] "GoogleKnowlege_Occupation"  
## [3] "Show"  
## [4] "Group"  
## [5] "Raw_Guest_List"
```


Renaming columns

To rename these columns, use `rename`. The basic syntax is:

```
## Generic code
rename(dataframe,
       new_column_name_1 = old_column_name_1,
       new_column_name_2 = old_column_name_2)
```

If you want to change column names in the saved object, be sure you reassign the object to be the output of `rename`.

Renaming columns

To rename columns in the `daily_show` data, then, use:

```
daily_show <- rename(daily_show,  
                     year = YEAR,  
                     job = GoogleKnowledge_Occupation,  
                     date = Show,  
                     category = Group,  
                     guest_name = Raw_Guest_List)  
  
head(daily_show, 3)
```

```
## # A tibble: 3 x 5  
##   year job          date    category guest_name  
##   <int> <chr>        <chr>   <chr>    <chr>  
## 1  1999 actor       1/11/~ Acting Michael J. ~  
## 2  1999 Comedian  1/12/~ Comedy Sandra Bern~  
## 3  1999 television ~ 1/13/~ Acting Tracey Ullm~
```

Renaming columns

As a quick check, what is the difference between these two calls?

1.

```
rename(daily_show,  
       year = YEAR,  
       job = GoogleKnowledge_Occupation,  
       date = Show,  
       category = Group,  
       guest_name = Raw_Guest_List)
```

2.

```
daily_show <- rename(daily_show,  
                    year = YEAR,  
                    job = GoogleKnowledge_Occupation,  
                    date = Show,  
                    category = Group,  
                    guest_name = Raw_Guest_List)
```

Selecting columns

Next, you may want to select only some columns of the dataframe. You can use `select` for this. The basic structure of this command is:

```
## Generic code  
select(dataframe, column_name_1, column_name_2, ...)
```

Where `column_name_1`, `column_name_2`, etc., are the names of the columns you want to keep.

Selecting columns

For example, to select all columns except year (since that information is already included in date), run:

```
select(daily_show, job, date, category, guest_name)
```

```
## # A tibble: 2,693 x 4
```

##	job	date	category	guest_name
##	<chr>	<chr>	<chr>	<chr>
## 1	actor	1/11/99	Acting	Michael J~
## 2	Comedian	1/12/99	Comedy	Sandra Be~
## 3	television actress	1/13/99	Acting	Tracey Ul~
## 4	film actress	1/14/99	Acting	Gillian A~
## 5	actor	1/18/99	Acting	David Ala~
## 6	actor	1/19/99	Acting	William B~
## 7	Singer-lyricist	1/20/99	Musician	Michael S~
## 8	model	1/21/99	Media	Carmen El~
## 9	actor	1/25/99	Acting	Matthew L~
## 10	stand-up comedian	1/26/99	Comedy	David Cro~

Selecting columns

As a reminder, we could have selected these columns using square bracket indexing, too:

```
daily_show[ , 2:5]
```

However, the `select` function will fit in nicely with other data-cleaning functions from “tidyverse” packages, plus the `select` function has some cool extra options, including:

- Selecting all columns that start with a certain pattern
- Selecting all columns that end with a certain pattern
- Selecting all columns that contain a certain pattern

Selecting columns

The `select` function also provides some time-saving tools. For example, in the last example, we wanted all the columns except one. Instead of writing out all the columns we want, we can use `-` with the columns we don't want to save time:

```
daily_show <- select(daily_show, -year)
head(daily_show, 3)
```

```
## # A tibble: 3 x 4
##   job          date    category guest_name
##   <chr>      <chr>    <chr>    <chr>
## 1 actor      1/11/99 Acting Michael J.~
## 2 Comedian   1/12/99 Comedy  Sandra Ber~
## 3 television actress 1/13/99 Acting  Tracey Ull~
```

Selecting columns

Another cool trick with `select` is that, if you want to keep several columns in a row, you can use a colon (`:`) with column names (rather than column position numbers) to select those columns:

```
daily_show <- select(daily_show, job:guest_name)
```

This call says that we want to select all columns from the one named “job” to the one named “guest_name”.

Add or change columns

You can change a column or add a new column using the `mutate` function. That function has the syntax:

```
# Generic code  
mutate(dataframe,  
        changed_column = function(changed_column),  
        new_column = function(other arguments))
```

- If you want to just **change** a column (in place), use its original name on the left of the equation.
- If you want to **add** a new column, use a new name on the left of the equation (this will be the name of the new column).

Add or change columns

For example, the job column in daily_show sometimes uses upper case and sometimes does not:

```
head(unique(daily_show$job), 10)
```

```
## [1] "actor" "Comedian"
## [3] "television actress" "film actress"
## [5] "Singer-lyricist" "model"
## [7] "stand-up comedian" "actress"
## [9] "comedian" "Singer-songwriter"
```

Add or change columns

We could use the `tolower` function to make all listings lowercase:

```
library("stringr")
daily_show <- mutate(daily_show, job = str_to_lower(job))
head(daily_show, 3)
```

```
## # A tibble: 3 x 4
##   job          date    category guest_name
##   <chr>      <chr>    <chr>    <chr>
## 1 actor      1/11/99 Acting Michael J.~
## 2 comedian   1/12/99 Comedy  Sandra Ber~
## 3 television 1/13/99 Acting  Tracey Ull~
```

We'll take a break now and do section 2.6.4 of the In-Course Exercise.

Dates in R

A common task when changing or adding columns is to change the class of some of the columns. This is especially common for dates, which will often be read in as a character vector when reading data into R.

Vector classes

Here are a few common vector classes in R:

Class	Example
character	"Chemistry", "Physics", "Mathematics"
numeric	10, 20, 30, 40
factor	Male [underlying number: 1], Female [2]
Date	"2010-01-01" [underlying number: 14,610]
logical	TRUE, FALSE

Vector classes

To find out the class of a vector, you can use `class()`:

```
class(daily_show$date)
```

```
## [1] "character"
```

Note: You can use `str` to get information on the classes of all columns in a dataframe. It's also printed at the top of output from `dplyr` functions.

lubridate package

In many cases you can use functions from the lubridate package to parse dates pretty easily.

For example, if you have a character string with the date in the order of *year-month-day*, you can use the `ymd` function from lubridate to convert the character string to the Date class. For example:

```
library("lubridate")  
my_date <- ymd("2008-10-13")  
class(my_date)
```

```
## [1] "Date"
```

lubridate package

The lubridate package has a number of functions for converting character strings into dates (or date-times). To decide which one to use, you just need to know the order of the elements of the date in the character string.

For example, here are some commonly-used lubridate functions:

lubridate function	Order of date elements
ymd	year-month-day
dmy	day-month-year
mdy_hm	month-day-year-hour-minute
ymd_hms	year-month-day-hour-minute-second

(Remember, you can use `vignette("lubridate")` and `?lubridate` to get help with the lubridate package.)

lubridate package

You will see dates represented in many different ways. For example, October might be included in data as “October”, “Oct”, or “10”. Further, the way the elements are separated can vary.

The functions in `lubridate` are pretty good at working with these different options intelligently:

```
mdy("10-31-2017")
```

```
## [1] "2017-10-31"
```

```
dmy("31 October 2017")
```

```
## [1] "2017-10-31"
```

Some more examples:

```
ymd_hms("2017/10/31--17:33:10")
```

```
## [1] "2017-10-31 17:33:10 UTC"
```

```
mdy_hm("Oct. 31, 2017 5:33PM", tz = "MST")
```

```
## [1] "2017-10-31 17:33:00 MST"
```

Converting to Date class

We can use the `mdy` function from `lubridate` to convert the date column in the `daily_show` dataset to a `Date` class:

```
daily_show <- mutate(daily_show, date = mdy(date))  
head(daily_show, 3)
```

```
## # A tibble: 3 x 4  
##   job          date          category guest_name  
##   <chr>        <date>        <chr>    <chr>  
## 1 actor      1999-01-11 Acting    Michael J. F~  
## 2 comedian  1999-01-12 Comedy    Sandra Bernh~  
## 3 television a~ 1999-01-13 Acting    Tracey Ullman
```

```
class(daily_show$date)
```

```
## [1] "Date"
```

Converting to Date class

Once you have an object in the Date class, you can do things like plot by date, calculate the range of dates, and calculate the total number of days the dataset covers:

```
range(daily_show$date)
```

```
## [1] "1999-01-11" "2015-08-05"
```

```
diff(range(daily_show$date))
```

```
## Time difference of 6050 days
```

lubridate package

The lubridate package also includes functions to pull out certain elements of a date. For example, we could use `wday` to create a new column with the weekday of each show:

```
daily_show <- mutate(daily_show, show_day = wday(date, label = T))  
head(select(daily_show, date, show_day), 3)
```

```
## # A tibble: 3 x 2  
##   date          show_day  
##   <date>        <ord>  
## 1 1999-01-11 Mon  
## 2 1999-01-12 Tue  
## 3 1999-01-13 Wed
```

Other functions in `lubridate` for pulling elements from a date include:

- `mday`: Day of the month
- `yday`: Day of the year
- `month`: Month
- `quarter`: Fiscal quarter
- `year`: Year

Filtering and logical operators

Slicing to certain rows

Last week, you learned how to use square bracket indexing to limit a dataframe to certain rows by row number:

```
daily_show[1:3, ]
```

The dplyr package has a function you can use to do this, called `slice`. That function has the syntax:

```
# Generic code  
slice(dataframe, starting_row:ending_row)
```

where `starting_row` is the row number of the first row you want to keep and `ending_row` is the row number of the last line you want to keep.

Slicing to certain rows

For example, to print the first three rows of the `daily_show` data, you can run:

```
slice(daily_show, 1:3)
```

```
## # A tibble: 3 x 5
##   job      date      category guest_name show_day
##   <chr>   <date>    <chr>    <chr>      <ord>
## 1 actor   1999-01-11 Acting   Michael J~ Mon
## 2 comedi~ 1999-01-12 Comedy   Sandra Be~ Tue
## 3 televi~ 1999-01-13 Acting   Tracey Ul~ Wed
```

Arranging rows

There is also a function, `arrange`, you can use to re-order the rows in a dataframe. The syntax for this function is:

```
# Generic code  
arrange(dataframe, column_to_order_by)
```

If you run this function to use a character vector to order, it will order the rows alphabetically by the values in that column. If you specify a numeric vector, it will order the rows by the numeric value.

Arranging rows

For example, we could reorder the `daily_show` data alphabetically by the values in the `category` column with the following call:

```
daily_show <- arrange(daily_show, category)
head(daily_show, 3)
```

```
## # A tibble: 3 x 5
##   job      date      category guest_name show_day
##   <chr>   <date>    <chr>    <chr>      <ord>
## 1 profe~ 2001-10-03 Academic Stephen S.~ Wed
## 2 profe~ 2001-12-03 Academic Nadine Str~ Mon
## 3 histo~ 2003-11-04 Academic Michael Be~ Tue
```

Arranging rows

If you want the ordering to be reversed (e.g., from “z” to “a” for character vectors, from higher to lower for numeric, latest to earliest for a Date), you can include the `desc` function.

For example, to reorder the `daily_show` data by descending date (latest to earliest), you can run:

```
daily_show <- arrange(daily_show, desc(date))  
head(daily_show, 3)
```

```
## # A tibble: 3 x 5  
##   job      date      category guest_name show_day  
##   <chr>   <date>    <chr>    <chr>      <ord>  
## 1 comedi~ 2015-08-05 Comedy   Louis C.K. Wed  
## 2 actor   2015-08-04 Acting    Denis Lea~ Tue  
## 3 stand~ 2015-08-03 Comedy   Amy Schum~ Mon
```

Filtering to certain rows

Next, you might want to filter the dataset down so that it only includes certain rows. You can use `filter` to do that. The syntax is:

```
## Generic code  
filter(dataframe, logical statement)
```

The `logical statement` gives the condition that a row must meet to be included in the output data frame. For example, you might want to pull:

- Rows from 2015
- Rows where the guest was an academic
- Rows where the job is not missing

Filtering to certain rows

For example, if you want to create a data frame that only includes guests who were scientists, you can run:

```
scientists <- filter(daily_show, category == "Science")  
head(scientists)
```

```
## # A tibble: 6 x 5  
##   job      date      category guest_name show_day  
##   <chr>  <date>    <chr>    <chr>      <ord>  
## 1 astro~ 2015-04-23 Science Neil deGra~ Thu  
## 2 surge~ 2014-10-06 Science Atul Gawan~ Mon  
## 3 astro~ 2013-09-04 Science Mario Livio Wed  
## 4 astro~ 2013-03-06 Science Neil deGra~ Wed  
## 5 prima~ 2012-04-16 Science Jane Gooda~ Mon  
## 6 astro~ 2012-02-27 Science Neil deGra~ Mon
```


Common logical operators in R

To build a logical statement to use in `filter`, you'll need to know some of R's logical operators:

Operator	Meaning	Example
<code>==</code>	equals	<code>category == "Acting"</code>
<code>!=</code>	does not equal	<code>category != "Comedy"</code>
<code>%in%</code>	is in	<code>category %in% c("Academic", "Science")</code>
<code>is.na()</code>	is NA	<code>is.na(job)</code>
<code>!is.na()</code>	is not NA	<code>!is.na(job)</code>
<code>&</code>	and	<code>year == 2015 & category == "Academic"</code>
<code> </code>	or	<code>year == 2015 category == "Academic"</code>

dplyr versus base R

Just so you know, all of these actions also have alternatives in base R:

dplyr	Base R equivalent
rename	Reassign colnames
select	Square bracket indexing
slice	Square bracket indexing
filter	subset
mutate	Use \$ to change / create columns

You will see these alternatives used in older code examples.

We'll take a break now and do section 2.6.5 of the In-Course Exercise.

Piping

Piping



Piping

If you look at the format of these dplyr functions, you'll notice that they all take a dataframe as their first argument:

Generic code

```
rename(dataframe,  
       new_column_name_1 = old_column_name_1,  
       new_column_name_2 = old_column_name_2)  
select(dataframe, column_name_1, column_name_2)  
slice(dataframe, starting_row:ending_row)  
arrange(dataframe, column_to_order_by)  
filter(dataframe, logical statement)  
mutate(dataframe,  
       changed_column = function(changed_column),  
       new_column = function(other arguments))
```

Piping

Classically, you would clean up a dataframe in R by reassigning the dataframe object at each step:

```
daily_show <- read_csv("../data/daily_show_guests.csv",  
                        skip = 4)  
daily_show <- rename(daily_show,  
                     job = GoogleKnowledge_Occupation,  
                     date = Show,  
                     category = Group,  
                     guest_name = Raw_Guest_List)  
daily_show <- select(daily_show, -YEAR)  
daily_show <- mutate(daily_show, job = str_to_lower(job))  
daily_show <- filter(daily_show, category == "Science")
```

“Piping” lets you clean this code up a bit. It can be used with any function that inputs a dataframe as its first argument. It “pipes” the dataframe created right before the pipe (`%>%`) into the function right after the pipe.

Piping

With piping, the same data cleaning looks like:

```
daily_show <- read_csv("../data/daily_show_guests.csv",  
                        skip = 4) %>%  
  rename(job = GoogleKnowledge_Occupation,  
         date = Show,  
         category = Group,  
         guest_name = Raw_Guest_List) %>%  
  select(-YEAR) %>%  
  mutate(job = str_to_lower(job)) %>%  
  filter(category == "Science")
```

Piping

Piping tip #1: As you are trying to figure out what “piped” code like this is doing, try highlighting from the start of the code through just part of the pipe and run that. For example, try highlighting and running just from `read_csv` through the before the `%>%` in the line with `select`, and see what that output looks like.

```
daily_show <- read_csv("../data/daily_show_guests.csv",  
                        skip = 4) %>%  
  rename(job = GoogleKnowledge_Occupation,  
         date = Show,  
         category = Group,  
         guest_name = Raw_Guest_List) %>%  
  select(-YEAR) %>%  
  mutate(job = str_to_lower(job)) %>%  
  filter(category == "Science")
```

Piping

Piping tip #2: When you are writing an R script that uses piping, first write it and make sure you have it right **without** assigning it to an R object (i.e., no `<-`). Often you'll use piping to clean up an object in R, but if you have to work on the piping code, you end up with different versions of the object, which will cause frustrations.

```
read_csv("../data/daily_show_guests.csv", skip = 4) %>%  
  rename(job = GoogleKnowlege_Occupation,  
         date = Show,  
         category = Group,  
         guest_name = Raw_Guest_List) %>%  
  select(-YEAR) %>%  
  mutate(job = str_to_lower(job)) %>%  
  filter(category == "Science")
```

Piping tip #3: There is a keyboard shortcut for the pipe symbol:

`Command-Shift-m`

We'll take a break now and do section 2.6.6 of the In-Course Exercise.